

```
In [1]: from pathlib import Path
        from typing import List

        import numpy as np
        import pandas as pd
        import torch
        import torch.nn as nn
        import matplotlib; matplotlib.use("tkAgg")
        import matplotlib.pyplot as plt
        from sklearn.metrics import (
            confusion_matrix,
            ConfusionMatrixDisplay,
            roc_curve,
            auc,
            RocCurveDisplay,
            precision_recall_curve,
            PrecisionRecallDisplay,
        )
        from sklearn.preprocessing import StandardScaler
```

```
In [2]: # A simple Transformer-based classifier for sequence data
        class TransformerClassifier(nn.Module):

            def __init__(
                self,
                input_size: int = 1,      # Number of features per time step
                seq_length: int = 12,     # Length of input sequences
                d_model: int = 64,        # Size of embedding / hidden dimension
                nhead: int = 4,          # Number of attention heads
                num_layers: int = 2,      # Number of Transformer layers
                dropout: float = 0.3,    # Dropout rate for regularisation
            ):
                super().__init__()

                # Project input features to model dimension
                self.input_projection = nn.Linear(input_size, d_model)

                # One encoder layer: self-attention + feed-forward
                enc_layer = nn.TransformerEncoderLayer(
                    d_model=d_model, nhead=nhead, dropout=dropout, batch_first=True
                )

                # Stack multiple encoder layers
                self.encoder = nn.TransformerEncoder(enc_layer, num_layers=num_layers)

                # Final classification head: transform to a single output
                self.classifier = nn.Sequential(
                    nn.Linear(d_model, 32),
                    nn.ReLU(),
                    nn.Dropout(dropout),
                    nn.Linear(32, 1)
                )

            def forward(self, x: torch.Tensor) -> torch.Tensor:
```

```

# project input to model dimension
x = self.input_projection(x)

# apply Transformer encoder
x = self.encoder(x)

# average across sequence (global mean pooling)
pooled = x.mean(dim=1)

# classify and apply sigmoid (output between 0 and 1)
return torch.sigmoid(self.classifier(pooled)).squeeze()

```

```

In [3]: # Wrap features (X) and Labels (y) into a DataLoader for batching
def _to_loader(X, y, batch_size: int, shuffle: bool = False):
    dataset = list(zip(X, y)) # Combine inputs and Labels
    return torch.utils.data.DataLoader(dataset, batch_size=batch_size, shuffle=shuf

# Train the model for one epoch (one full pass through the data)
def train_one_epoch(model, loader, criterion, optim):
    model.train() # Set model to training mode
    running = 0.0 # Track total Loss

    for xb, yb in loader: # Loop through batches
        optim.zero_grad() # Reset gradients
        loss = criterion(model(xb), yb) # Compute Loss
        loss.backward() # Backpropagate
        optim.step() # Update weights
        running += loss.item() * xb.size(0) # Accumulate batch Loss

    return running / len(loader.dataset) # Return average Loss

# Evaluate the model's accuracy
def evaluate(model, loader):
    model.eval() # Set model to evaluation mode
    correct = 0

    with torch.no_grad(): # No gradient tracking needed
        for xb, yb in loader:
            preds = (model(xb) >= 0.5).float() # Predict and round to 0 or 1
            correct += (preds == yb).sum().item() # Count correct predictions

    return correct / len(loader.dataset) # Return accuracy

```

```

In [4]: # Manually performs cross-validation using custom folds in your DataFrame
def cross_validate_manual(
    df: pd.DataFrame,
    feature_cols: List[str], # List of columns to use as features
    epochs: int = 20,
    batch_size: int = 64,
    lr: float = 3e-4, # Learning rate
    weight_decay: float = 1e-3, # Weight decay for regularization
):
    results = [] # Store best validation accuracy for each fold

```

```

# Loop through each unique fold number
for fold in sorted(df["Fold"].unique()):
    print(f"\n— Fold {fold + 1} / {df['Fold'].nunique()} —————")

    # Split into training and validation sets
    train_df = df[df["Fold"] != fold]
    val_df = df[df["Fold"] == fold]

    # Normalize features and reshape for model input
    scaler = StandardScaler()
    X_train = scaler.fit_transform(train_df[feature_cols]).reshape(-1, 12, 1)
    X_val = scaler.transform(val_df[feature_cols]).reshape(-1, 12, 1)

    # Convert to PyTorch tensors
    y_train = train_df["Label"].values.astype(np.float32)
    y_val = val_df["Label"].values.astype(np.float32)
    X_train_t = torch.tensor(X_train, dtype=torch.float32)
    y_train_t = torch.tensor(y_train)
    X_val_t = torch.tensor(X_val, dtype=torch.float32)
    y_val_t = torch.tensor(y_val)

    # Create DataLoaders
    train_loader = _to_loader(X_train_t, y_train_t, batch_size, shuffle=True)
    val_loader = _to_loader(X_val_t, y_val_t, batch_size)

    # Initialize model, loss, and optimizer
    model = TransformerClassifier()
    criterion = nn.BCELoss() # Binary classification loss
    optim = torch.optim.AdamW(model.parameters(), lr=lr, weight_decay=weight_de

    best_acc = 0.0
    bad_epochs = 0
    patience = 5 # Early stopping if no improvement for 'patience' epochs

    # Training loop
    for epoch in range(1, epochs + 1):
        loss = train_one_epoch(model, train_loader, criterion, optim)
        val_acc = evaluate(model, val_loader)
        print(f"Epoch {epoch:02}/{epochs} - loss: {loss:.4f} - val acc: {val_ac

        # Save best model based on validation accuracy
        if val_acc > best_acc:
            best_acc, bad_epochs = val_acc, 0
            best_state = model.state_dict()
        else:
            bad_epochs += 1
            if bad_epochs == patience:
                print("Early stopping")
                break

    # Save best accuracy for the fold
    results.append(best_acc)

    # Save best model checkpoint
    Path("checkpoints").mkdir(exist_ok=True)

```

```

        torch.save(best_state, f"checkpoints/fold_{fold}.pt")
        print(f"Best acc fold {fold + 1}: {best_acc:.4f}")

    # Summary of results
    print("\n==== Validation Accuracy Summary =====")
    for i, acc in enumerate(results, 1):
        print(f"Fold {i}: {acc:.4f}")
    print(f"Mean Accuracy: {np.mean(results):.4f}")
    print(f"Standard Deviation: {np.std(results):.4f}")

    return results

```

```

In [5]: # Trains on the full dataset and shows evaluation visualisations
def visual_evaluation_full(df: pd.DataFrame, feature_cols: List[str]):
    # Scale features and reshape for the model
    scaler = StandardScaler()
    X = scaler.fit_transform(df[feature_cols]).reshape(-1, 12, 1)
    y = df["Label"].values.astype(np.float32)

    # Convert to PyTorch tensors
    X_t = torch.tensor(X, dtype=torch.float32)
    y_t = torch.tensor(y)

    # Create DataLoader
    loader = _to_loader(X_t, y_t, batch_size=64, shuffle=True)

    # Initialize model, loss function, and optimizer
    model = TransformerClassifier()
    criterion = nn.BCELoss()
    optim = torch.optim.AdamW(model.parameters(), lr=3e-4, weight_decay=1e-3)

    # Train for a fixed number of epochs
    for _ in range(15):
        train_one_epoch(model, loader, criterion, optim)

    # Get model predictions
    model.eval()
    with torch.no_grad():
        y_scores = model(X_t).cpu().numpy().ravel() # Raw output scores
        y_pred = (y_scores > 0.5).astype(int) # Convert scores to binary predictions

    # Confusion Matrix
    cm = confusion_matrix(y, y_pred)
    ConfusionMatrixDisplay(confusion_matrix=cm).plot(cmap="Blues")
    plt.title("Confusion Matrix")
    plt.savefig("confusion_matrix.png", dpi=300, bbox_inches="tight")
    plt.show(block=True)

    # ROC Curve
    fpr, tpr, _ = roc_curve(y, y_scores)
    roc_auc = auc(fpr, tpr)
    RocCurveDisplay(fpr=fpr, tpr=tpr, roc_auc=roc_auc).plot()
    plt.title(f"ROC Curve (AUC = {roc_auc:.4f})")
    plt.savefig("roc_curve.png", dpi=300, bbox_inches="tight")
    plt.show(block=True)

```

```

# Precision-Recall Curve
precision, recall, _ = precision_recall_curve(y, y_scores)
PrecisionRecallDisplay(precision=precision, recall=recall).plot()
plt.title("Precision-Recall Curve")
plt.savefig("pr_curve.png", dpi=300, bbox_inches="tight")
plt.show(block=True)

```

```

In [6]: import time
import psutil
import os

if __name__ == "__main__":
    process = psutil.Process(os.getpid())

    # === Resource Monitoring Start ===
    overall_start_time = time.time()
    overall_start_ram = process.memory_info().rss / 1024 / 1024 # in MB
    overall_start_cpu = psutil.cpu_percent(interval=1)

    # Load dataset
    df = pd.read_csv("D:\\Coding Projects\\Detection-of-SYN-Flood-Attacks-Using-Mac")
    feature_cols = df.columns.difference(["Label", "Fold"]).tolist()[:12]

    # Run cross-validation on Transformer
    results = cross_validate_manual(df, feature_cols)

    # === Resource Monitoring End ===
    overall_end_time = time.time()
    overall_end_ram = process.memory_info().rss / 1024 / 1024 # in MB
    overall_end_cpu = psutil.cpu_percent(interval=1)

    # === Summary ===
    print("\n Overall Training Stats ")
    print(f"Total Training Time: {overall_end_time - overall_start_time:.2f} second")
    print(f"Total RAM Usage Increase: {overall_end_ram - overall_start_ram:.2f} MB")
    print(f"CPU Usage (at final check): {overall_end_cpu}%")

```

— Fold 1 / 5 —
Epoch 01/20 - loss: 0.2557 - val acc: 0.9969
Epoch 02/20 - loss: 0.0478 - val acc: 0.9969
Epoch 03/20 - loss: 0.0388 - val acc: 0.9969
Epoch 04/20 - loss: 0.0360 - val acc: 0.9969
Epoch 05/20 - loss: 0.0300 - val acc: 0.9969
Epoch 06/20 - loss: 0.0302 - val acc: 0.9964
Early stopping
Best acc fold 1: 0.9969

— Fold 2 / 5 —
Epoch 01/20 - loss: 0.2690 - val acc: 0.9958
Epoch 02/20 - loss: 0.0498 - val acc: 0.9958
Epoch 03/20 - loss: 0.0324 - val acc: 0.9958
Epoch 04/20 - loss: 0.0262 - val acc: 0.9958
Epoch 05/20 - loss: 0.0227 - val acc: 0.9958
Epoch 06/20 - loss: 0.0244 - val acc: 0.9870
Early stopping
Best acc fold 2: 0.9958

— Fold 3 / 5 —
Epoch 01/20 - loss: 0.1876 - val acc: 0.9927
Epoch 02/20 - loss: 0.0420 - val acc: 0.9927
Epoch 03/20 - loss: 0.0336 - val acc: 0.9927
Epoch 04/20 - loss: 0.0317 - val acc: 0.9927
Epoch 05/20 - loss: 0.0256 - val acc: 0.9927
Epoch 06/20 - loss: 0.0254 - val acc: 0.9927
Early stopping
Best acc fold 3: 0.9927

— Fold 4 / 5 —
Epoch 01/20 - loss: 0.2297 - val acc: 0.9932
Epoch 02/20 - loss: 0.0332 - val acc: 0.9958
Epoch 03/20 - loss: 0.0222 - val acc: 0.9964
Epoch 04/20 - loss: 0.0227 - val acc: 0.9917
Epoch 05/20 - loss: 0.0214 - val acc: 0.9927
Epoch 06/20 - loss: 0.0160 - val acc: 0.9964
Epoch 07/20 - loss: 0.0163 - val acc: 0.9969
Epoch 08/20 - loss: 0.0134 - val acc: 0.9969
Epoch 09/20 - loss: 0.0145 - val acc: 0.9969
Epoch 10/20 - loss: 0.0154 - val acc: 0.9969
Epoch 11/20 - loss: 0.0134 - val acc: 0.9969
Epoch 12/20 - loss: 0.0120 - val acc: 0.9969
Early stopping
Best acc fold 4: 0.9969

— Fold 5 / 5 —
Epoch 01/20 - loss: 0.2440 - val acc: 0.9932
Epoch 02/20 - loss: 0.0389 - val acc: 0.9932
Epoch 03/20 - loss: 0.0280 - val acc: 0.9932
Epoch 04/20 - loss: 0.0273 - val acc: 0.9917
Epoch 05/20 - loss: 0.0240 - val acc: 0.9932
Epoch 06/20 - loss: 0.0190 - val acc: 0.9932
Early stopping
Best acc fold 5: 0.9932

===== Validation Accuracy Summary =====

Fold 1: 0.9969

Fold 2: 0.9958

Fold 3: 0.9927

Fold 4: 0.9969

Fold 5: 0.9932

Mean Accuracy: 0.9951

Standard Deviation: 0.0018

Overall Training Stats

Total Training Time: 153.87 seconds

Total RAM Usage Increase: 187.30 MB

CPU Usage (at final check): 7.7%