

```
In [ ]: from pathlib import Path
from typing import List
import numpy as np
import pandas as pd
import torch
import torch.nn as nn
import matplotlib; matplotlib.use("tkAgg")
import matplotlib.pyplot as plt
from sklearn.metrics import (
    confusion_matrix,
    ConfusionMatrixDisplay,
    roc_curve,
    auc,
    RocCurveDisplay,
    precision_recall_curve,
    PrecisionRecallDisplay,
)
from sklearn.preprocessing import StandardScaler
```

```
In [ ]: # simple Transformer-based classifier for sequence data
class TransformerClassifier(nn.Module):

    def __init__(
        self,
        input_size: int = 1,      # number of features per time step, 1 as we have
        seq_length: int = 12,     # length of input sequences, covers all the 12 fe
        d_model: int = 64,        # size of embedding vector, kept relatively small
        nhead: int = 4,           # number of attention heads, kept small for speed
        num_layers: int = 2,      # number of Transformer layers, kept small for sp
        dropout: float = 0.3,     # dropout rate for regularisation, set 0.3 to red
    ):
        super().__init__()

        # project input features to model dimension
        self.input_projection = nn.Linear(input_size, d_model)

        # one encoder layer: self-attention + feed-forward
        enc_layer = nn.TransformerEncoderLayer(
            d_model=d_model, nhead=nhead, dropout=dropout, batch_first=True
        )

        # stack the two encoder layers
        self.encoder = nn.TransformerEncoder(enc_layer, num_layers=num_layers)

        # final classification head
        self.classifier = nn.Sequential(
            nn.Linear(d_model, 32),
            nn.ReLU(),
            nn.Dropout(dropout),
            nn.Linear(32, 1)
        )

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        x = self.input_projection(x)
```

```

        x = self.encoder(x)
        pooled = x.mean(dim=1) # mean pooling over the sequence length as suggested
        return torch.sigmoid(self.classifier(pooled)).squeeze()

```

```

In [ ]: # wrap the features (X) and labels (y) into a DataLoader for batching
def _to_loader(X, y, batch_size: int, shuffle: bool = False):
    dataset = list(zip(X, y))
    return torch.utils.data.DataLoader(dataset, batch_size=batch_size, shuffle=shuf

# train the model for one epoch at a time
def train_one_epoch(model, loader, criterion, optim):
    model.train()
    running = 0.0

    for xb, yb in loader:
        optim.zero_grad()
        loss = criterion(model(xb), yb)
        loss.backward()
        optim.step()
        running += loss.item() * xb.size(0)

    return running / len(loader.dataset)

# Evaluate the model's accuracy
def evaluate(model, loader):
    model.eval()
    correct = 0

    with torch.no_grad():
        for xb, yb in loader:
            preds = (model(xb) >= 0.5).float()
            correct += (preds == yb).sum().item()
    return correct / len(loader.dataset)

```

```

In [ ]: # manually perform the cross-validation using custom k-folds in the DataFrame
def cross_validate_manual(
    Syn_df: pd.DataFrame,
    feature_columns: List[str],
    epochs: int = 20,
    batch_size: int = 64,
    lr: float = 3e-4,
    weight_decay: float = 1e-3,
):
    results = [] # store best validation accuracy for each fold

    # Loop through each unique fold number
    for fold in sorted(Syn_df["Fold"].unique()):
        print(f"\n— Fold {fold + 1} / {Syn_df['Fold'].nunique()} —————")

        # split into both training and validation sets
        train_df = Syn_df[Syn_df["Fold"] != fold]
        validate_df = Syn_df[Syn_df["Fold"] == fold]

        # normalize features here and then reshape for the model input

```

```

standard_scaler = StandardScaler()
X_train = standard_scaler.fit_transform(train_df[feature_columns]).reshape(
X_validate = standard_scaler.transform(validate_df[feature_columns]).reshap

# convert to PyTorch tensors
y_train = train_df["Label"].values.astype(np.float32)
y_validate = validate_df["Label"].values.astype(np.float32)
X_train_ten = torch.tensor(X_train, dtype=torch.float32)
y_train_ten = torch.tensor(y_train)
X_val_ten = torch.tensor(X_validate, dtype=torch.float32)
y_val_ten = torch.tensor(y_validate)

# create the dataLoaders
train_loader = _to_loader(X_train_ten, y_train_ten, batch_size, shuffle=True)
val_loader = _to_loader(X_val_ten, y_val_ten, batch_size)

# initialize model, loss, and optimizer
transformer_model = TransformerClassifier()
criterion = nn.BCELoss() # binary classification loss
optim = torch.optim.AdamW(transformer_model.parameters(), lr=lr, weight_dec

best_accuracy = 0.0
inaccurate_epochs = 0
patience = 5 # early stopping if no improvement for 'patience' epochs

# training loop
for epoch in range(1, epochs + 1):
    loss = train_one_epoch(transformer_model, train_loader, criterion, opti
    val_acc = evaluate(transformer_model, val_loader)
    print(f"Epoch {epoch:02}/{epochs} - loss: {loss:.4f} - val acc: {val_ac

    # save the best model based on validation accuracy
    if val_acc > best_accuracy:
        best_accuracy, inaccurate_epochs = val_acc, 0
        best_state = transformer_model.state_dict()
    else:
        inaccurate_epochs += 1
        if inaccurate_epochs == patience:
            print("Early stopping")
            break

# save the best accuracy for the fold
results.append(best_accuracy)

# save the best model checkpoint
Path("checkpoints").mkdir(exist_ok=True)
torch.save(best_state, f"checkpoints/fold_{fold}.pt")
print(f"Best acc fold {fold + 1}: {best_accuracy:.4f}")

# summary of the final results
print("\n===== Validation Accuracy Summary =====")
for i, acc in enumerate(results, 1):
    print(f"Fold {i}: {acc:.4f}")
print(f"Mean Accuracy: {np.mean(results):.4f}")
print(f"Standard Deviation: {np.std(results):.4f}")

```

```
return results
```

```
In [ ]: # trains on the full dataset and shows evaluation visualisations
def visual_evaluation_full(Syn_df: pd.DataFrame, feature_columns: List[str]):
    # Scale features and reshape for the model
    standard_scaler = StandardScaler()
    X = standard_scaler.fit_transform(Syn_df[feature_columns]).reshape(-1, 12, 1)
    y = Syn_df["Label"].values.astype(np.float32)

    # convert to PyTorch tensors
    X_t = torch.tensor(X, dtype=torch.float32)
    y_t = torch.tensor(y)

    # create DataLoader
    loader = _to_loader(X_t, y_t, batch_size=64, shuffle=True)

    # initialize model, loss function, and optimizer
    transformer_model = TransformerClassifier()
    criterion = nn.BCELoss()
    optim = torch.optim.AdamW(transformer_model.parameters(), lr=3e-4, weight_decay

    # train for a fixed number of epochs
    for _ in range(15):
        train_one_epoch(transformer_model, loader, criterion, optim)

    # get model predictions
    transformer_model.eval()
    with torch.no_grad():
        y_scores = transformer_model(X_t).cpu().numpy().ravel()
    y_pred = (y_scores > 0.5).astype(int)

    # confusion Matrix
    cm = confusion_matrix(y, y_pred)
    ConfusionMatrixDisplay(confusion_matrix=cm).plot(cmap="Blues")
    plt.title("Confusion Matrix")
    plt.savefig("confusion_matrix.png", dpi=300, bbox_inches="tight")
    plt.show(block=True)

    # ROC Curve
    fpr, tpr, _ = roc_curve(y, y_scores)
    roc_auc = auc(fpr, tpr)
    RocCurveDisplay(fpr=fpr, tpr=tpr, roc_auc=roc_auc).plot()
    plt.title(f"ROC Curve (AUC = {roc_auc:.4f})")
    plt.savefig("roc_curve.png", dpi=300, bbox_inches="tight")
    plt.show(block=True)

    # Precision-Recall Curve
    precision, recall, _ = precision_recall_curve(y, y_scores)
    PrecisionRecallDisplay(precision=precision, recall=recall).plot()
    plt.title("Precision-Recall Curve")
    plt.savefig("pr_curve.png", dpi=300, bbox_inches="tight")
    plt.show(block=True)
```

```

In [ ]: import time
import psutil
import os

if __name__ == "__main__":
    process = psutil.Process(os.getpid())

    # resource monitoring start point
    overall_start_time = time.time()
    overall_start_ram = process.memory_info().rss / 1024 / 1024 # in MB
    overall_start_cpu = psutil.cpu_percent(interval=1)

    # Load dataset
    Syn_df = pd.read_csv("D:\\Coding Projects\\Detection-of-SYN-Flood-Attacks-Using
feature_columns = Syn_df.columns.difference(["Label", "Fold"]).tolist()[12]

    # run cross-validation on Transformer
    results = cross_validate_manual(Syn_df, feature_columns)

    # end resource monitoring
    overall_end_time = time.time()
    overall_end_ram = process.memory_info().rss / 1024 / 1024 # in MB
    overall_end_cpu = psutil.cpu_percent(interval=1)

    # summary of trainig stats
    print("\n Overall Training Stats ")
    print(f"Total Training Time: {overall_end_time - overall_start_time:.2f} second")
    print(f"Total RAM Usage Increase: {overall_end_ram - overall_start_ram:.2f} MB")
    print(f"CPU Usage (at final check): {overall_end_cpu}%")

```

— Fold 1 / 5 —  
Epoch 01/20 - loss: 0.2557 - val acc: 0.9969  
Epoch 02/20 - loss: 0.0478 - val acc: 0.9969  
Epoch 03/20 - loss: 0.0388 - val acc: 0.9969  
Epoch 04/20 - loss: 0.0360 - val acc: 0.9969  
Epoch 05/20 - loss: 0.0300 - val acc: 0.9969  
Epoch 06/20 - loss: 0.0302 - val acc: 0.9964  
Early stopping  
Best acc fold 1: 0.9969

— Fold 2 / 5 —  
Epoch 01/20 - loss: 0.2690 - val acc: 0.9958  
Epoch 02/20 - loss: 0.0498 - val acc: 0.9958  
Epoch 03/20 - loss: 0.0324 - val acc: 0.9958  
Epoch 04/20 - loss: 0.0262 - val acc: 0.9958  
Epoch 05/20 - loss: 0.0227 - val acc: 0.9958  
Epoch 06/20 - loss: 0.0244 - val acc: 0.9870  
Early stopping  
Best acc fold 2: 0.9958

— Fold 3 / 5 —  
Epoch 01/20 - loss: 0.1876 - val acc: 0.9927  
Epoch 02/20 - loss: 0.0420 - val acc: 0.9927  
Epoch 03/20 - loss: 0.0336 - val acc: 0.9927  
Epoch 04/20 - loss: 0.0317 - val acc: 0.9927  
Epoch 05/20 - loss: 0.0256 - val acc: 0.9927  
Epoch 06/20 - loss: 0.0254 - val acc: 0.9927  
Early stopping  
Best acc fold 3: 0.9927

— Fold 4 / 5 —  
Epoch 01/20 - loss: 0.2297 - val acc: 0.9932  
Epoch 02/20 - loss: 0.0332 - val acc: 0.9958  
Epoch 03/20 - loss: 0.0222 - val acc: 0.9964  
Epoch 04/20 - loss: 0.0227 - val acc: 0.9917  
Epoch 05/20 - loss: 0.0214 - val acc: 0.9927  
Epoch 06/20 - loss: 0.0160 - val acc: 0.9964  
Epoch 07/20 - loss: 0.0163 - val acc: 0.9969  
Epoch 08/20 - loss: 0.0134 - val acc: 0.9969  
Epoch 09/20 - loss: 0.0145 - val acc: 0.9969  
Epoch 10/20 - loss: 0.0154 - val acc: 0.9969  
Epoch 11/20 - loss: 0.0134 - val acc: 0.9969  
Epoch 12/20 - loss: 0.0120 - val acc: 0.9969  
Early stopping  
Best acc fold 4: 0.9969

— Fold 5 / 5 —  
Epoch 01/20 - loss: 0.2440 - val acc: 0.9932  
Epoch 02/20 - loss: 0.0389 - val acc: 0.9932  
Epoch 03/20 - loss: 0.0280 - val acc: 0.9932  
Epoch 04/20 - loss: 0.0273 - val acc: 0.9917  
Epoch 05/20 - loss: 0.0240 - val acc: 0.9932  
Epoch 06/20 - loss: 0.0190 - val acc: 0.9932  
Early stopping  
Best acc fold 5: 0.9932

===== Validation Accuracy Summary =====

Fold 1: 0.9969  
Fold 2: 0.9958  
Fold 3: 0.9927  
Fold 4: 0.9969  
Fold 5: 0.9932  
Mean Accuracy: 0.9951  
Standard Deviation: 0.0018

Overall Training Stats

Total Training Time: 153.87 seconds  
Total RAM Usage Increase: 187.30 MB  
CPU Usage (at final check): 7.7%

## saving the model as PDF

```
In [ ]: !jupyter nbconvert --to webpdf "d:\\Coding Projects\\Detection-of-SYN-Flood-Attacks
```

```
[NbConvertApp] Converting notebook d:\\Coding Projects\\Detection-of-SYN-Flood-Attac  
ks-Using-Machine-Learning-and-Deep-Learning-Techniques-with-Feature-Base\\Taulant Ma  
tarova\\Transformer updated.ipynb to webpdf  
[NbConvertApp] Building PDF  
[NbConvertApp] PDF successfully created  
[NbConvertApp] Writing 131790 bytes to d:\\Coding Projects\\Detection-of-SYN-Flood-Att  
acks-Using-Machine-Learning-and-Deep-Learning-Techniques-with-Feature-Base\\Taulant M  
atarova\\Transformer updated.pdf
```