

```

from pathlib import Path
from typing import List
import numpy as np
import pandas as pd
import torch
import torch.nn as nn
import matplotlib
import scipy
import matplotlib.pyplot as plt
from sklearn.metrics import (
    confusion_matrix,
    ConfusionMatrixDisplay,
    roc_curve,
    auc,
    RocCurveDisplay,
    precision_recall_curve,
    PrecisionRecallDisplay,
)
from sklearn.preprocessing import StandardScaler

# simple Transformer-based classifier for sequence data
class TransformerClassifier(nn.Module):

    def __init__(
        self,
        input_size: int = 1,      # number of features per time step,
        1 as we have 12 features so not much to process
        seq_length: int = 12,    # length of input sequences, covers
        all the 12 features
        d_model: int = 64,       # size of embedding vector, kept
        relatively small for speed
        nhead: int = 4,          # number of attention heads, kept
        small for speed
        num_layers: int = 2,     # number of Transformer layers, kept
        small for speed
        dropout: float = 0.3,    # dropout rate for regularisation,
        set 0.3 to reduce overfitting
    ):
        super().__init__()

        # project input features to model dimension
        self.input_projection = nn.Linear(input_size, d_model)

        # one encoder layer: self-attention + feed-forward
        enc_layer = nn.TransformerEncoderLayer(
            d_model=d_model, nhead=nhead, dropout=dropout,
            batch_first=True
        )

        # stack the two encoder layers

```

```

        self.encoder = nn.TransformerEncoder(enc_layer,
num_layers=num_layers)

    # final classification head
    self.classifier = nn.Sequential(
        nn.Linear(d_model, 32),
        nn.ReLU(),
        nn.Dropout(dropout),
        nn.Linear(32, 1)
    )

    def forward(self, x: torch.Tensor) -> torch.Tensor:
        x = self.input_projection(x)
        x = self.encoder(x)
        pooled = x.mean(dim=1) # mean pooling over the sequence length
as suggested by ChatGPT
        return torch.sigmoid(self.classifier(pooled)).squeeze()

# wrap the features (X) and labels (y) into a DataLoader for batching
    def _to_loader(X, y, batch_size: int, shuffle: bool = False):
        dataset = list(zip(X, y))
        return torch.utils.data.DataLoader(dataset, batch_size=batch_size,
shuffle=shuffle)

# train the model for one epoch at a time
    def train_one_epoch(model, loader, criterion, optim):
        model.train()
        running = 0.0

        for xb, yb in loader:
            optim.zero_grad()
            loss = criterion(model(xb), yb)
            loss.backward()
            optim.step()
            running += loss.item() * xb.size(0)

        return running / len(loader.dataset)

# Evaluate the model's accuracy
    def evaluate(model, loader):
        model.eval()
        correct = 0

        with torch.no_grad():
            for xb, yb in loader:
                preds = (model(xb) >= 0.5).float()
                correct += (preds == yb).sum().item()
        return correct / len(loader.dataset)

```

```

device = "cuda" if torch.cuda.is_available() else "cpu"
%matplotlib inline
# manually perform the cross-validation using custom k-folds in the DataFrame
def cross_validate_manual(
    Syn_df: pd.DataFrame,
    feature_columns: List[str],
    epochs: int = 20,
    batch_size: int = 64,
    lr: float = 3e-4,
    weight_decay: float = 1e-3,
):
    results = [] # store best validation accuracy for each fold
    oof_true: List[int] = []
    oof_pred: List[int] = []
    oof_score: List[float] = []

    # loop through each unique fold number
    for fold in sorted(Syn_df["Fold"].unique()):
        print(f"\n— Fold {fold + 1} / {Syn_df['Fold'].nunique()}")

        # split into both training and validation sets
        train_df = Syn_df[Syn_df["Fold"] != fold]
        validate_df = Syn_df[Syn_df["Fold"] == fold]

        # normalize features here and then reshape for the model input
        standard_scaler = StandardScaler()
        X_train =
standard_scaler.fit_transform(train_df[feature_columns]).reshape(-1,
12, 1)
        X_validate =
standard_scaler.transform(validate_df[feature_columns]).reshape(-1,
12, 1)

        # convert to PyTorch tensors
        y_train = train_df["Label"].values.astype(np.float32)
        y_validate = validate_df["Label"].values.astype(np.float32)
        X_train_ten = torch.tensor(X_train, dtype=torch.float32,
device=device)
        y_train_ten = torch.tensor(y_train, device=device)
        X_val_ten = torch.tensor(X_validate, dtype=torch.float32,
device=device)
        y_val_ten = torch.tensor(y_validate, device=device)

        # create the dataLoaders
        train_loader = _to_loader(X_train_ten, y_train_ten,
batch_size, shuffle=True)
        val_loader = _to_loader(X_val_ten, y_val_ten, batch_size)

```

```

    # initialize model, loss, and optimizer
    transformer_model = TransformerClassifier().to(device) # move
model to GPU if available
    criterion = nn.BCELoss() # binary classification loss
    optim = torch.optim.AdamW(transformer_model.parameters(),
lr=lr, weight_decay=weight_decay)

    best_accuracy = 0.0
    inaccurate_epochs = 0
    patience = 5 # early stopping if no improvement for
'patience' epochs

    # training loop
    for epoch in range(1, epochs + 1):
        loss = train_one_epoch(transformer_model, train_loader,
criterion, optim)
        val_acc = evaluate(transformer_model, val_loader)
        print(f"Epoch {epoch:02}/{epochs} - loss: {loss:.4f} - val
acc: {val_acc:.4f}")

        # save the best model based on validation accuracy
        if val_acc > best_accuracy:
            best_accuracy, inaccurate_epochs = val_acc, 0
            best_state = transformer_model.state_dict()
        else:
            inaccurate_epochs += 1
            if inaccurate_epochs == patience:
                print("Early stopping")
                break

        # get the model predictions for the validation set
        transformer_model.eval()
        with torch.no_grad():
            y_prob =
transformer_model(X_val_ten).cpu().numpy().ravel()

            y_pred = (y_prob > 0.5).astype(int) # threshold to 0/1
            oof_true.extend(y_validate.tolist()) # true labels
            oof_pred.extend(y_pred.tolist()) # predicted labels
            oof_score.extend(y_prob.tolist())

        # save the best accuracy for the fold
        results.append(best_accuracy)

        # save the best model checkpoint
        Path("checkpoints").mkdir(exist_ok=True)
        torch.save(best_state, f"checkpoints/fold_{fold}.pt")
        print(f"Best acc fold {fold + 1}: {best_accuracy:.4f}")

    # summary of the final results
    print("\n===== Validation Accuracy Summary =====")

```

```

for i, acc in enumerate(results, 1):
    print(f"Fold {i}: {acc:.4f}")
print(f"Mean Accuracy: {np.mean(results):.4f}")
print(f"Standard Deviation: {np.std(results):.4f}")

cm = confusion_matrix(oof_true, oof_pred)
ConfusionMatrixDisplay(confusion_matrix=cm).plot(cmap='Blues')
plt.title('Transformer Confusion Matrix')
plt.show()

fpr, tpr, _ = roc_curve(oof_true, oof_score)
RocCurveDisplay(fpr=fpr, tpr=tpr, roc_auc=auc(fpr, tpr)).plot()
plt.title('Transformer ROC')
plt.show()

prec, rec, _ = precision_recall_curve(oof_true, oof_score)
PrecisionRecallDisplay(precision=prec, recall=rec,
average_precision=pr_auc).plot()
plt.title(f'Transformer PR (AUC = {pr_auc:.3f})')
plt.show()

return results

import psutil, os, time, numpy as np, pandas as pd
process = psutil.Process(os.getpid())

# resource monitoring start point
overall_start_time = time.time()
overall_start_ram = process.memory_info().rss / 1024 / 1024 # in MB
overall_start_cpu = psutil.cpu_percent(interval=1)

# Load dataset
Syn_df = pd.read_csv("D:\\Coding Projects\\Detection-of-SYN-Flood-
Attacks-Using-Machine-Learning-and-Deep-Learning-Techniques-with-
Feature-Base\\Data\\K5_Dataset.csv")
feature_columns = Syn_df.columns.difference(["Label",
"Fold"]).tolist()[12]

# run cross-validation on Transformer
results = cross_validate_manual(Syn_df, feature_columns)

# end resource monitoring
overall_end_time = time.time()
overall_end_ram = process.memory_info().rss / 1024 / 1024 # in MB
overall_end_cpu = psutil.cpu_percent(interval=1)

# summary of training stats
print("\n Overall Training Stats ")

```

```

print(f"Total Training Time: {overall_end_time -
overall_start_time:.2f} seconds")
print(f"Total RAM Usage Increase: {overall_end_ram -
overall_start_ram:.2f} MB")
print(f"CPU Usage (at final check): {overall_end_cpu}%")
print("Per-fold accuracies      :", [f"{x:.4f}" for x in results])
print(f"Mean Accuracy          : {np.mean(results):.4f}")
print(f"Std-Dev                 : {np.std(results):.4f}")

# save the notebook as a web PDF
os.getcwd()
!jupyter nbconvert --to webpdf "d:\\Coding Projects\\Detection-of-SYN-
Flood-Attacks-Using-Machine-Learning-and-Deep-Learning-Techniques-
with-Feature-Base\\Taulant Matarova\\Transformer_model_2.ipynb"

```

— Fold 1 / 5 —

```

Epoch 01/20 - loss: 0.2189 - val acc: 0.9969
Epoch 02/20 - loss: 0.0452 - val acc: 0.9969
Epoch 03/20 - loss: 0.0372 - val acc: 0.9969
Epoch 04/20 - loss: 0.0336 - val acc: 0.9974
Epoch 05/20 - loss: 0.0283 - val acc: 0.9953
Epoch 06/20 - loss: 0.0260 - val acc: 0.9964
Epoch 07/20 - loss: 0.0237 - val acc: 0.9964
Epoch 08/20 - loss: 0.0238 - val acc: 0.9901
Epoch 09/20 - loss: 0.0244 - val acc: 0.9969
Early stopping
Best acc fold 1: 0.9974

```

— Fold 2 / 5 —

```

Epoch 01/20 - loss: 0.2515 - val acc: 0.9958
Epoch 02/20 - loss: 0.0450 - val acc: 0.9958
Epoch 03/20 - loss: 0.0371 - val acc: 0.9958
Epoch 04/20 - loss: 0.0286 - val acc: 0.9953
Epoch 05/20 - loss: 0.0268 - val acc: 0.9917
Epoch 06/20 - loss: 0.0223 - val acc: 0.9922
Early stopping
Best acc fold 2: 0.9958

```

— Fold 3 / 5 —

```

Epoch 01/20 - loss: 0.2029 - val acc: 0.9927
Epoch 02/20 - loss: 0.0402 - val acc: 0.9927
Epoch 03/20 - loss: 0.0315 - val acc: 0.9927
Epoch 04/20 - loss: 0.0276 - val acc: 0.9927
Epoch 05/20 - loss: 0.0252 - val acc: 0.9885
Epoch 06/20 - loss: 0.0242 - val acc: 0.9901
Early stopping
Best acc fold 3: 0.9927


```

— Fold 4 / 5 —

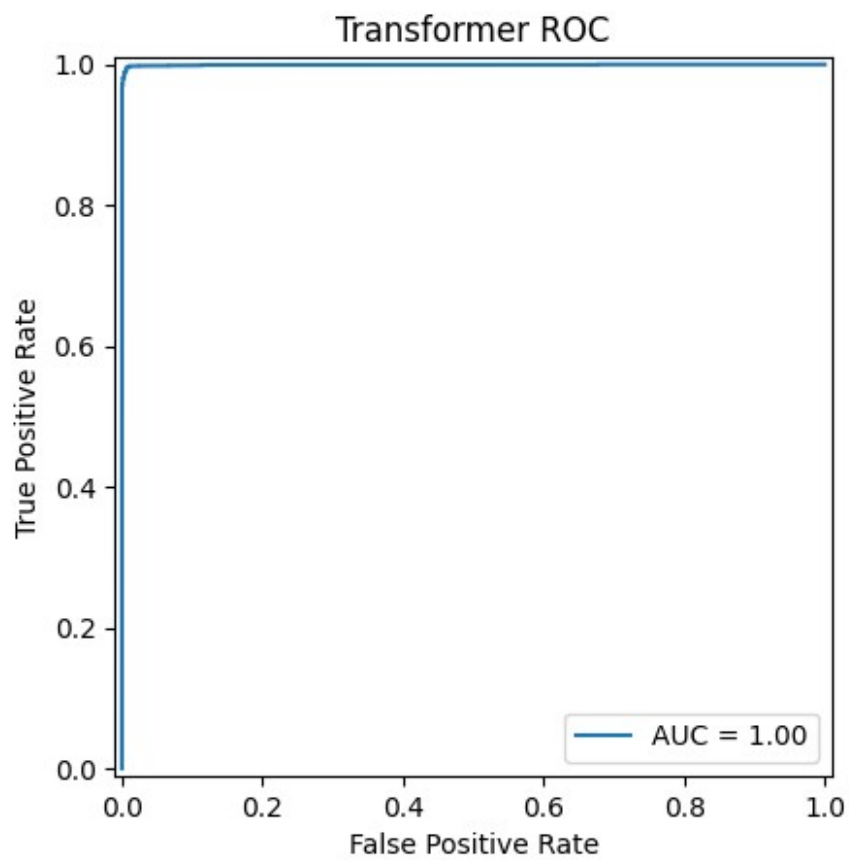
Epoch 01/20 - loss: 0.2448 - val acc: 0.9958  
Epoch 02/20 - loss: 0.0411 - val acc: 0.9958  
Epoch 03/20 - loss: 0.0348 - val acc: 0.9958  
Epoch 04/20 - loss: 0.0270 - val acc: 0.9958  
Epoch 05/20 - loss: 0.0246 - val acc: 0.9958  
Epoch 06/20 - loss: 0.0232 - val acc: 0.9958  
Early stopping  
Best acc fold 4: 0.9958

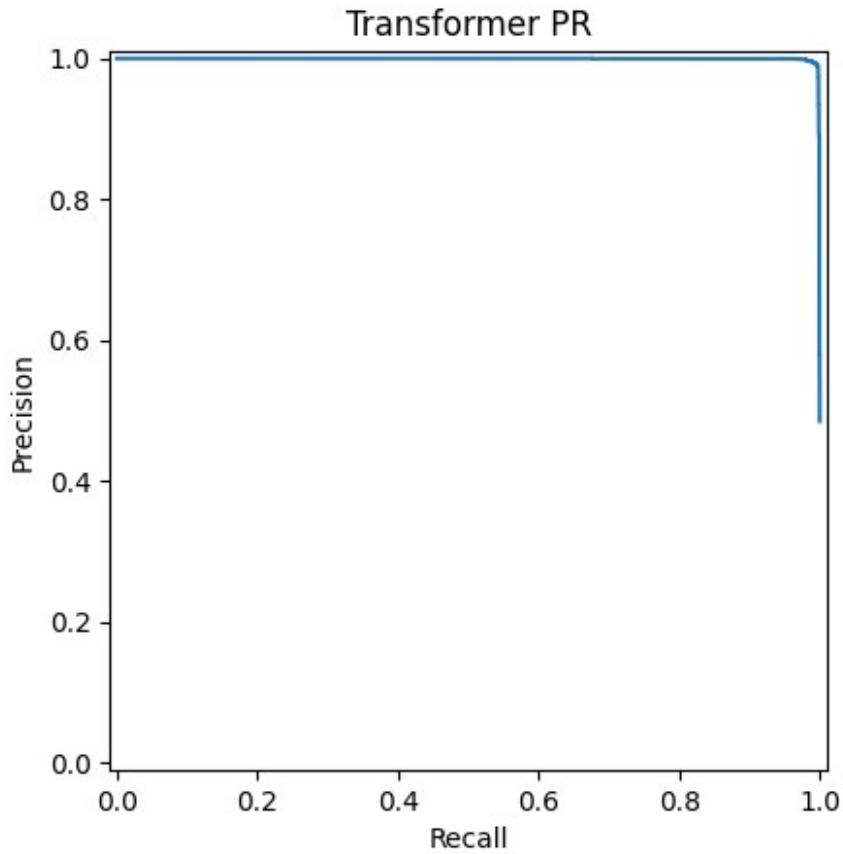
— Fold 5 / 5 —  
Epoch 01/20 - loss: 0.2592 - val acc: 0.9932  
Epoch 02/20 - loss: 0.0378 - val acc: 0.9932  
Epoch 03/20 - loss: 0.0288 - val acc: 0.9932  
Epoch 04/20 - loss: 0.0260 - val acc: 0.9932  
Epoch 05/20 - loss: 0.0234 - val acc: 0.9932  
Epoch 06/20 - loss: 0.0250 - val acc: 0.9932  
Early stopping  
Best acc fold 5: 0.9932

===== Validation Accuracy Summary =====  
Fold 1: 0.9974  
Fold 2: 0.9958  
Fold 3: 0.9927  
Fold 4: 0.9958  
Fold 5: 0.9932  
Mean Accuracy: 0.9950  
Standard Deviation: 0.0018









```
Overall Training Stats
Total Training Time: 141.70 seconds
Total RAM Usage Increase: 356.96 MB
CPU Usage (at final check): 4.4%
Per-fold accuracies      : ['0.9974', '0.9958', '0.9927', '0.9958',
                             '0.9932']
Mean Accuracy            : 0.9950
Std-Dev                  : 0.0018

[NbConvertApp] Converting notebook d:\\Coding Projects\\Detection-of-
SYN-Flood-Attacks-Using-Machine-Learning-and-Deep-Learning-Techniques-
with-Feature-Base\\Taulant Matarova\\Transformer_model_2.ipynb to
webpdf
[NbConvertApp] Building PDF
[NbConvertApp] PDF successfully created
[NbConvertApp] Writing 119433 bytes to d:\\Coding Projects\\Detection-
of-SYN-Flood-Attacks-Using-Machine-Learning-and-Deep-Learning-
Techniques-with-Feature-Base\\Taulant Matarova\\Transformer_model_2.pdf
```