

```
In [ ]: from pathlib import Path
from typing import List
import numpy as np
import pandas as pd
import torch
import torch.nn as nn
import matplotlib; matplotlib.use("tkAgg")
import matplotlib.pyplot as plt
from sklearn.metrics import (
    confusion_matrix,
    ConfusionMatrixDisplay,
    roc_curve,
    auc,
    RocCurveDisplay,
    precision_recall_curve,
    PrecisionRecallDisplay,
    average_precision_score
)
from sklearn.preprocessing import StandardScaler
```

```
In [17]: # simple Transformer-based classifier for sequence data
class TransformerClassifier(nn.Module):

    def __init__(
        self,
        input_size: int = 1,      # number of features per time step, 1 as we have
        seq_length: int = 12,     # length of input sequences, covers all the 12 fe
        d_model: int = 64,        # size of embedding vector, kept relatively small
        nhead: int = 4,           # number of attention heads, kept small for speed
        num_layers: int = 2,      # number of Transformer layers, kept small for sp
        dropout: float = 0.3,     # dropout rate for regularisation, set 0.3 to red
    ):
        super().__init__()

        # project input features to model dimension
        self.input_projection = nn.Linear(input_size, d_model)

        # one encoder layer: self-attention + feed-forward
        enc_layer = nn.TransformerEncoderLayer(
            d_model=d_model, nhead=nhead, dropout=dropout, batch_first=True
        )

        # stack the two encoder layers
        self.encoder = nn.TransformerEncoder(enc_layer, num_layers=num_layers)

        # final classification head
        self.classifier = nn.Sequential(
            nn.Linear(d_model, 32),
            nn.ReLU(),
            nn.Dropout(dropout),
            nn.Linear(32, 1)
        )

    def forward(self, x: torch.Tensor) -> torch.Tensor:
```

```

        x = self.input_projection(x)
        x = self.encoder(x)
        pooled = x.mean(dim=1) # mean pooling over the sequence length as suggested
        return torch.sigmoid(self.classifier(pooled)).squeeze()

```

```

In [18]: # wrap the features (X) and labels (y) into a DataLoader for batching
def _to_loader(X, y, batch_size: int, shuffle: bool = False):
    dataset = list(zip(X, y))
    return torch.utils.data.DataLoader(dataset, batch_size=batch_size, shuffle=shuffle)

# train the model for one epoch at a time
def train_one_epoch(model, loader, criterion, optim):
    model.train()
    running = 0.0

    for xb, yb in loader:
        optim.zero_grad()
        loss = criterion(model(xb), yb)
        loss.backward()
        optim.step()
        running += loss.item() * xb.size(0)

    return running / len(loader.dataset)

# Evaluate the model's accuracy
def evaluate(model, loader):
    model.eval()
    correct = 0

    with torch.no_grad():
        for xb, yb in loader:
            preds = (model(xb) >= 0.5).float()
            correct += (preds == yb).sum().item()
    return correct / len(loader.dataset)

```

```

In [ ]: device = "cuda" if torch.cuda.is_available() else "cpu"

# manually perform the cross-validation using custom k-folds in the DataFrame
def cross_validate_manual(
    Syn_df: pd.DataFrame,
    feature_columns: List[str],
    epochs: int = 20,
    batch_size: int = 64,
    lr: float = 3e-4,
    weight_decay: float = 1e-3,
):
    results = [] # store best validation accuracy for each fold
    oof_true, oof_score = [], [] # collectors for out-of-fold predictions

    # Loop through each unique fold number
    for fold in sorted(Syn_df["Fold"].unique()):
        print(f"\n— Fold {fold + 1} / {Syn_df['Fold'].nunique()} —————")
        # split into both training and validation sets

```

```

train_df = Syn_df[Syn_df["Fold"] != fold]
validate_df = Syn_df[Syn_df["Fold"] == fold]

# normalize features here and then reshape for the model input
standard_scaler = StandardScaler()
X_train = standard_scaler.fit_transform(train_df[feature_columns]).reshape(
X_validate = standard_scaler.transform(validate_df[feature_columns]).reshape(

# convert to PyTorch tensors
y_train = train_df["Label"].values.astype(np.float32)
y_validate = validate_df["Label"].values.astype(np.float32)
X_train_ten = torch.tensor(X_train, dtype=torch.float32, device=device)
y_train_ten = torch.tensor(y_train, device=device)
X_val_ten = torch.tensor(X_validate, dtype=torch.float32, device=device)
y_val_ten = torch.tensor(y_validate, device=device)

# create the dataLoaders
train_loader = _to_loader(X_train_ten, y_train_ten, batch_size, shuffle=True)
val_loader = _to_loader(X_val_ten, y_val_ten, batch_size)

# initialize model, loss, and optimizer
transformer_model = TransformerClassifier().to(device) # move model to GPU
criterion = nn.BCELoss() # binary classification loss
optim = torch.optim.AdamW(transformer_model.parameters(), lr=lr, weight_dec

best_accuracy = 0.0
inaccurate_epochs = 0
patience = 5 # early stopping if no improvement for 'patience' epochs

# training loop
for epoch in range(1, epochs + 1):
    loss = train_one_epoch(transformer_model, train_loader, criterion, optim)
    val_acc = evaluate(transformer_model, val_loader)
    print(f"Epoch {epoch:02}/{epochs} - loss: {loss:.4f} - val acc: {val_acc:.4f}")

    # save the best model based on validation accuracy
    if val_acc > best_accuracy:
        best_accuracy, inaccurate_epochs = val_acc, 0
        best_state = transformer_model.state_dict()
    else:
        inaccurate_epochs += 1
        if inaccurate_epochs == patience:
            print("Early stopping")
            break

# get the model predictions for the validation set
transformer_model.eval()
with torch.no_grad():
    y_prob = transformer_model(X_val_ten).cpu().numpy().ravel()
oof_true.extend(y_validate.tolist())
oof_score.extend(y_prob.tolist())

# save the best accuracy for the fold
results.append(best_accuracy)

# save the best model checkpoint

```

```

    Path("checkpoints").mkdir(exist_ok=True)
    torch.save(best_state, f"checkpoints/fold_{fold}.pt")
    print(f"Best acc fold {fold + 1}: {best_accuracy:.4f}")

# summary of the final results
    print("\n==== Validation Accuracy Summary =====")
    for i, acc in enumerate(results, 1):
        print(f"Fold {i}: {acc:.4f}")
    print(f"Mean Accuracy: {np.mean(results):.4f}")
    print(f"Standard Deviation: {np.std(results):.4f}")
    y_bin = (np.array(oof_score) > 0.5).astype(int)

    cm = confusion_matrix(oof_true, y_bin)
    ConfusionMatrixDisplay(confusion_matrix=cm).plot(cmap='Blues')
    plt.title('Transformer Confusion Matrix')
    plt.show()
    print(cm)

    fpr, tpr, _ = roc_curve(oof_true, oof_score)
    RocCurveDisplay(fpr=fpr, tpr=tpr, roc_auc=auc(fpr, tpr)).plot()
    plt.title('Transformer ROC')
    plt.show()
    print(fpr, tpr, auc(fpr, tpr))

    prec, rec, _ = precision_recall_curve(oof_true, oof_score)
    PrecisionRecallDisplay(precision=prec, recall=rec).plot()
    plt.title('Transformer PR')
    plt.show()
    print("PR AUC:", average_precision_score(oof_true, oof_score))

    return results

```

Using: cpu

```

In [ ]: import time
import psutil
import os

process = psutil.Process(os.getpid())

# resource monitoring start point
overall_start_time = time.time()
overall_start_ram = process.memory_info().rss / 1024 / 1024 # in MB
overall_start_cpu = psutil.cpu_percent(interval=1)

# Load dataset
Syn_df = pd.read_csv("D:\\Coding Projects\\Detection-of-SYN-Flood-Attacks-Using-Mac
feature_columns = Syn_df.columns.difference(["Label", "Fold"]).tolist()[:12]

# run cross-validation on Transformer
results = cross_validate_manual(Syn_df, feature_columns)

# end resource monitoring
overall_end_time = time.time()
overall_end_ram = process.memory_info().rss / 1024 / 1024 # in MB

```

```
overall_end_cpu = psutil.cpu_percent(interval=1)

# summary of training stats
print("\n Overall Training Stats ")
print(f"Total Training Time: {overall_end_time - overall_start_time:.2f} seconds")
print(f"Total RAM Usage Increase: {overall_end_ram - overall_start_ram:.2f} MB")
print(f"CPU Usage (at final check): {overall_end_cpu}%")
print("Per-fold accuracies      :", [f"{x:.4f}" for x in results])
print(f"Mean Accuracy           : {np.mean(results):.4f}")
print(f"Std-Dev                  : {np.std(results):.4f}")
```

— Fold 1 / 5 —  
Epoch 01/20 - loss: 0.2156 - val acc: 0.9969  
Epoch 02/20 - loss: 0.0411 - val acc: 0.9969  
Epoch 03/20 - loss: 0.0341 - val acc: 0.9969  
Epoch 04/20 - loss: 0.0278 - val acc: 0.9969  
Epoch 05/20 - loss: 0.0281 - val acc: 0.9969  
Epoch 06/20 - loss: 0.0254 - val acc: 0.9969  
Early stopping  
Best acc fold 1: 0.9969

— Fold 2 / 5 —  
Epoch 01/20 - loss: 0.2257 - val acc: 0.9958  
Epoch 02/20 - loss: 0.0403 - val acc: 0.9958  
Epoch 03/20 - loss: 0.0318 - val acc: 0.9958  
Epoch 04/20 - loss: 0.0291 - val acc: 0.9958  
Epoch 05/20 - loss: 0.0260 - val acc: 0.9958  
Epoch 06/20 - loss: 0.0230 - val acc: 0.9880  
Early stopping  
Best acc fold 2: 0.9958

— Fold 3 / 5 —  
Epoch 01/20 - loss: 0.2193 - val acc: 0.9927  
Epoch 02/20 - loss: 0.0411 - val acc: 0.9927  
Epoch 03/20 - loss: 0.0310 - val acc: 0.9927  
Epoch 04/20 - loss: 0.0264 - val acc: 0.9927  
Epoch 05/20 - loss: 0.0255 - val acc: 0.9927  
Epoch 06/20 - loss: 0.0164 - val acc: 0.9964  
Epoch 07/20 - loss: 0.0154 - val acc: 0.9964  
Epoch 08/20 - loss: 0.0170 - val acc: 0.9964  
Epoch 09/20 - loss: 0.0152 - val acc: 0.9964  
Epoch 10/20 - loss: 0.0170 - val acc: 0.9964  
Epoch 11/20 - loss: 0.0136 - val acc: 0.9964  
Early stopping  
Best acc fold 3: 0.9964

— Fold 4 / 5 —  
Epoch 01/20 - loss: 0.2373 - val acc: 0.9922  
Epoch 02/20 - loss: 0.0372 - val acc: 0.9958  
Epoch 03/20 - loss: 0.0232 - val acc: 0.9943  
Epoch 04/20 - loss: 0.0220 - val acc: 0.9906  
Epoch 05/20 - loss: 0.0276 - val acc: 0.9927  
Epoch 06/20 - loss: 0.0177 - val acc: 0.9906  
Epoch 07/20 - loss: 0.0175 - val acc: 0.9932  
Early stopping  
Best acc fold 4: 0.9958

— Fold 5 / 5 —  
Epoch 01/20 - loss: 0.2155 - val acc: 0.9932  
Epoch 02/20 - loss: 0.0369 - val acc: 0.9854  
Epoch 03/20 - loss: 0.0300 - val acc: 0.9932  
Epoch 04/20 - loss: 0.0255 - val acc: 0.9932  
Epoch 05/20 - loss: 0.0242 - val acc: 0.9932  
Epoch 06/20 - loss: 0.0230 - val acc: 0.9932  
Early stopping  
Best acc fold 5: 0.9932

===== Validation Accuracy Summary =====

Fold 1: 0.9969  
Fold 2: 0.9958  
Fold 3: 0.9964  
Fold 4: 0.9958  
Fold 5: 0.9932  
Mean Accuracy: 0.9956  
Standard Deviation: 0.0013

Overall Training Stats  
Total Training Time: 152.20 seconds  
Total RAM Usage Increase: 23.27 MB  
CPU Usage (at final check): 6.6%

## saving the model as PDF

```
In [21]: import os  
os.getcwd()
```

```
Out[21]: 'd:\\Coding Projects\\Detection-of-SYN-Flood-Attacks-Using-Machine-Learning-and-De  
ep-Learning-Techniques-with-Feature-Base\\Taulant Matarova'
```

```
In [22]: !jupyter nbconvert --to webpdf "d:\\Coding Projects\\Detection-of-SYN-Flood-Attacks
```

```
[NbConvertApp] Converting notebook d:\\Coding Projects\\Detection-of-SYN-Flood-Attac  
ks-Using-Machine-Learning-and-Deep-Learning-Techniques-with-Feature-Base\\Taulant Ma  
tarova\\Transformer_model_final.ipynb to webpdf  
[NbConvertApp] Building PDF  
[NbConvertApp] PDF successfully created  
[NbConvertApp] Writing 136222 bytes to d:\\Coding Projects\\Detection-of-SYN-Flood-Att  
acks-Using-Machine-Learning-and-Deep-Learning-Techniques-with-Feature-Base\\Taulant M  
atarova\\Transformer_model_final.pdf
```